

Modelling the Semantics of Adjectives in the Ontology-Lexicon Interface

John P. McCrae

Affiliation / Address line 1

Affiliation / Address line 2

Affiliation / Address line 3

email@domain

Francesca Quattri

Affiliation / Address line 1

Affiliation / Address line 2

Affiliation / Address line 3

email@domain

Christina Unger, Philipp Cimiano

Affiliation / Address line 1

Affiliation / Address line 2

Affiliation / Address line 3

email@domain

Abstract

The treatment of the semantics of adjectives is notoriously challenging, and we consider this problem in the context of the so called ontology-lexicon interface which attempts to capture the semantics of words by reference to an ontology in description logics or some other, typically first-order, logical formalism. The use of first order logic (hence also description logics), while effective for nouns and verbs, breaks down in the case of adjectives. We argue that this is primarily due to a lack of logical expressivity in the underlying ontology languages. In particular, beyond the straightforward *intersective adjectives*, there exist *gradable adjectives*, requiring fuzzy or non-monotonic semantics, as well as *operator adjectives*, requiring second-order logic for modelling. We consider how we can extend the ontology-lexicon interface as realized by extant models such as *lemon* in the face of the issues mentioned above, in particular those arising in the context of modelling the ontological semantics of adjectives. We show how more complex logical formalisms **that are** required to capture the ontological semantics of adjectives can be backward engineered into OWL-based modelling by means of pseudo-classes. We discuss the implications of this modelling in the context of an application **to ontology-based question answering...of an application (ontology-based question answering)..**

1 Introduction

Ontology-lexicon models, such as *lemon* (Lexicon Model for Ontologies) (McCrae et al., 2012), have become an important model for handling a number of tasks in natural language processing. In particular, such ontology-lexica are built around the separation of a lexical layer describing how a word or phrase acts syntactically and morphologically, and a semantic layer describing how the meaning of a word is expressed in a formal logical model, such as OWL (Web Ontology Language) (Deborah L. McGuinness and others, 2004). It has been shown that this principle known as *semantics by reference* (Buitelaar, 2010) is an effective model that can support the task of developing question answering systems (Unger and Cimiano, 2011) and natural language generation (Cimiano et al., 2013) over backends based on Semantic Web data models. The Pythia system, which builds on the *lemon* formalism to declaratively capture the lexicon-ontology interface, for example, has been instantiated to the case of answering questions from DBpedia (Unger and Cimiano, 2011) However, as has been shown by the Question Answering over Linked Data (Lopez et al., 2013, QALD) benchmarking campaigns, there are many questions that can be asked over this database that require a deeper representation of the semantics of words, adjectives in particular. For example, questions such as (1a) require understanding of the semantics of ‘high’ in a manner that goes beyond the expressivity of OWL. The answer given in the QALD dataset for this question is shown in (1b). In particular, the interpretation of this question involves the understanding of how the word ‘high’ relates to the property `dbo:elevation`, including ordering and subset selection operations, **and how to express and requires expressing** this semantics in a formal manner.

1. (a) What is the highest mountain in Australia?

```
(b) SELECT DISTINCT ?uri WHERE {
    ?uri rdf:type dbo:Mountain .
    ?uri dbo:locatedInArea res:Australia .
    ?uri dbo:elevation ?elevation .
} ORDER BY DESC(?elevation) LIMIT 1
```

It has been claimed that first-order logic and thus by extension description logics, such as OWL, “fail decidedly when it comes to adjectives” (Bankston, 2003). In fact, we largely agree that the semantics of many adjectives are difficult or impossible to describe in first-order logic. However, from the point of view of the ontology-lexicon interface, the logical expressivity of the ontology is not a limiting factor. In fact, due to the separation of the lexical and ontology layers in a model such as *lemon*, it is possible to express the meaning of words without worrying about the formalism used in the ontology. To this extent, we will first demonstrate that adjectives are in general a case where the use of description logics (DL) breaks down, and for which more sophisticated logical formalisms must be applied. We then consider to what extent this can be handled in the context of the ontology-lexicon, and introduce pseudo-classes, that is OWL classes with annotations, which we use to express the semantics of adjectives in a manner that would allow reasoning with fuzzy, high-order models. To this extent, we base our models on the previously introduced design patterns (M^cCrae and Unger, 2014) for modeling ontology-lexica. Finally, we show how these semantics can be helpful in practical applications of question answering over the DBpedia knowledge base.

2 Classification of adjectives

There are a number of classifications of adjectives. First we will start with the most fundamental distinction between *attributive* versus *predicative* usage, that is the use of adjectives in noun phrases (“*X* is a *A N*”) versus as objects of the copula (“*X* is *A*”). It should be noted that there are many adjectives for which only predicative or attributive usage is allowed, as shown in 3a and 3.

2. (a) Mary is a technical engineer.
- (b) *Mary is technical.
3. (a) The baby is awake.
- (b) *The awake baby.

One of the principle classifications of the semantics of adjectives (for example (Partee, 2003; Bouillon and Viegas, 1999; Morzycki, 2013b)) is based on the meaning of adjective noun compounds relative to the meaning of the words by themselves. This classification is as follows (where \Rightarrow denotes entailment).

Intersective (X is a *A N* $\Rightarrow X$ is *A* $\wedge X$ is a *N*) Such adjectives work as if they were another noun and indicate that the compound noun phrase is a member of both the class of the noun and the class of the adjective. For example, in the phrase “Belgian violinist” it refers to a person in the class intersection $Belgian \sqcap Violinist(X)$, and hence we can infer that a “Belgian violinist” is a subclass of a “Belgian”. Furthermore, we could conclude that if the same person were a surgeon, they would also be a “Belgian Surgeon”.

Subjective (X is a *A N* $\Rightarrow X$ is a *N*, but X is a *A N* $\not\Rightarrow X$ is *A*) Such adjectives acquire their specific meaning in combination with the noun they modify. For example, a “skilful violinist” is certainly in the class $Violinist(X)$ but the described person is ‘skilful as a violinist’, but not skilful in general, e.g. as a surgeon.

Privative (X is a *A N* $\not\Rightarrow X$ is a *N*) These adjectives modify the meaning of a noun phrase to create a noun phrase that is potentially incompatible with the original meaning. For example, a “fake gun” is not a member of the class of guns.

Another important distinction is whether adjectives are *gradable*, i.e. whether a comparative or superlative statement with these adjectives makes sense. For example, adjectives such as ‘big’ or ‘tall’ can express relationships such as ‘*X* is bigger than *Y*’. However it is not possible to say one individual is

‘more former’. Most gradable adjectives are subsective, for example ‘a big mouse’ is not ‘a big animal’ (Morzycki, 2013a).

Finally, we consider *operator* or *property-modifying* adjectives, which can be considered to be the same as privative adjectives, but in this case are understood as operators that change some property in the qualia structure of the class. For example, we may express the adjective ‘former’ in lambda calculus as a function that takes a class C as input and returns the class of entities that were a member of C to some prior time point t (Partee, 2003):

$$\lambda C[\lambda x \exists t C(x, t) \cap t < \text{now}]$$

Such adjectives have not only a difference in semantic meaning but can also frequently have syntactic impact, for example in adjective ordering restrictions, as they may be reordered with only semantic impact (Teodorescu, 2006), e.g.,

4. (a) A big red car.
(b) ?A red big car.
5. (a) A famous former actor.
(b) A former famous actor.

One further case that is important to distinguish is that of *relational* adjectives which have a meaning that expresses a relationship between two individuals or events, for example:

6. He is related to her.
7. She is similar to her brother.
8. This is useful for something.

3 Representation of adjectives in the ontology-lexicon interface

In general it is assumed that adjectives form frames with exactly one argument except for extra arguments given by adjuncts, typically prepositional phrases. Most adjectives are thus associated with a predicative frame, which much like the standard noun predicate frame (X is a N) is stereotyped in English as:

$$X \text{ is } A$$

For attributive, usage we associate this with a frame, which is stereotyped as, where the $N?$ argument is not semantically bound, but instead obtained by syntactic unification with a noun predicate frame.

$$X \text{ is } A \ N?$$

As such, when we encounter the attributive usage of an adjective such as in 9, we understand this as the realization of two frames, given in 10.

9. Juan is a Spanish researcher.
10. (a) Juan is a researcher.
(b) Juan is a Spanish ?.

Note we do not provide modeling for adjectives where the meaning is unique for a particular noun phrase, such as ‘polar bear’, which we would capture as a normal noun phrase with meaning *ursus maritimus*.

3.1 Intersective adjectives

Intersective adjectives are the most straightforward class as in many cases they can be modelled essentially as a noun or verb (e.g. deverbal adjectives such as ‘broken’). Intersective adjectives take one argument and can thus be modelled as unary predicates in first-order logic or classes in OWL as described by McCrae et al. (McCrae and Unger, 2014). For practical modeling examples we will use the *lemon* model as it is the most prominent implementation of the ontology-lexicon interface.

The primary mechanism of modelling the syntax-semantics interface in the context of *lemon* is by means of assigning a *frame* as a *syntactic behaviour* of an entry and giving it *syntactic arguments*, which can then be linked to the *lexical sense*, which stands in proxy for a true semantic frame in the ontology.

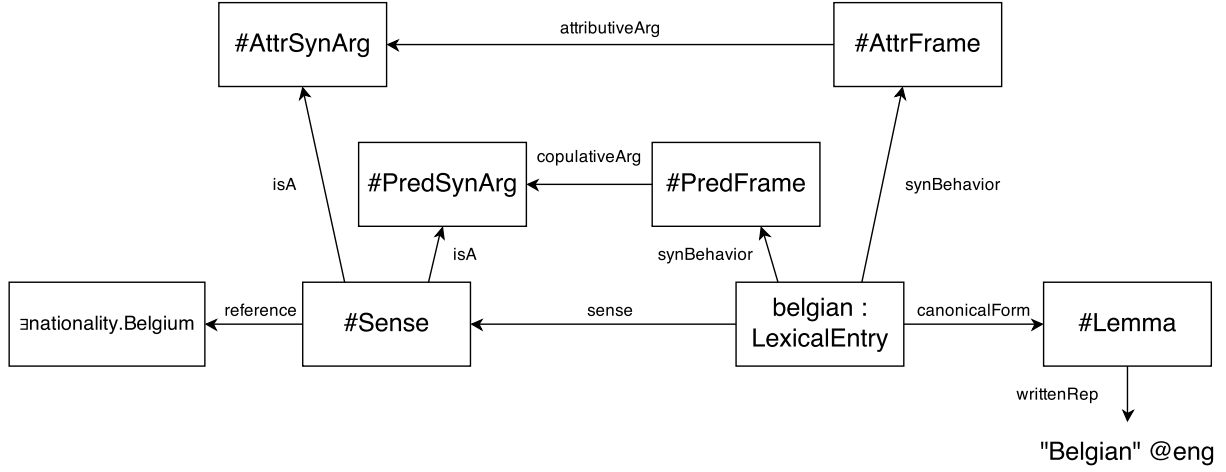


Figure 1: Modelling of an intersective adjective ‘Belgian’ in *lemon*

For example, the modelling of an adjective such as ‘Belgian’ can be achieved as follows (depicted in Figure 1)¹.

```
lexicon:belgian a lemon:LexicalEntry ;
  lemon:canonicalForm belgian:Lemma ;
  lemon:synBehavior    belgian:AttrFrame ,
                      belgian:PredFrame ;
  lemon:sense          belgian:Sense .

belgian:Lemma lemon:writtenRep "Belgian"@eng .

belgian:AttrFrame lexinfo:attributiveArg belgian:AttrSynArg .
belgian:PredFrame lexinfo:copulativeArg  belgian:PredSynArg .

belgian:sense lemon:reference [ a owl:Restriction ;
                               owl:onProperty dbpedia:nationality ;
                               owl:hasValue dbpedia:Belgium ] ;
  lemon:isA belgian:AttrSynArg , belgian:PredSynArg .
```

Note, that here we use the external vocabulary LexInfo (Cimiano et al., 2011) to define the meaning of the arguments of the frame as the *attributive argument*, corresponding to the frame stereotype ‘ X is A N ?’ and the *copulative argument* for the frame stereotype ‘ X is A ’. Furthermore, the class of Belgians is not named in our reference ontology DBpedia, so we introduce an anonymous class with the axiomatization, $\exists \text{nationality.Belgium}$. It is in fact common that the referent of an adjective is not named in an ontology and as such we tend to model denominal adjectives as classes of the form $\exists \text{prop.Value}$, where *Value* is the reference of the noun from which the adjective is derived. This modelling is so common that it has already been encoded as two design patterns, called *IntersectiveObjectPropertyAdjective* and *IntersectiveDatatypePropertyAdjective*, see McCrae and Unger (2014). Similarly, most deverbal adjectives refer to an event, and as such a common modelling is of the form $\exists \text{theme}^{-1}.EventClass$, for example ‘vandalized’ may be $\exists \text{theme}^{-1}.VandalismEvent$.

3.2 Gradable adjectives and relevant observables

Gradable adjectives have a number of properties, which differentiate them from intersective adjectives:

- They occur in comparative constructions, in English with either ‘-er’ or ‘more’ (Kennedy and McNally, 1999):3), e.g. ‘smaller’ and ‘more frequent’, as opposed to intersectives such as ‘*less geological’ and ‘*more wooden’.

¹We assume that the namespaces are defined for the lexicon as `lexicon`, e.g., `http://www.example.org/lexicon` and for the entry, e.g., `belgian` is `http://www.example.org/lexicon/belgian#`, other namespaces are assumed to be as usual.

- Gradable adjectives can be defined as ‘scalar’, since their value can ideally be measured on a scale of set degrees, which is dependent on the object of the adjective (‘big’ can be defined based on the number of stories of a ‘building’ or the height of a ‘person’).
- They have a context-dependent truth-conditional variability, meaning that their positive form is understood in relations to the class of the object modified by the adjective. In other words, an ‘expensive watch’ has a different price scale to an ‘expensive bottle of water’.
- They are frequently *fuzzy* (or *vague*) (Kennedy, 2007).
- There may be a minimum or maximum of the adjective’s scale, which can be determined by, for example, whether they can be modified by ‘completely’.

As such we define gradable adjectives relative to a particular property, these adjectives are also called ‘observable’ (Bennett, 2006)², that is it is natural to say that ‘big’ refers to ‘size’, however it is clear then that ‘small’ also refers to ‘size’ and as antonyms they cannot both refer to the same ontological concept. As such, we introduce the concept of *covariance* and *contravariance*, which refers to whether the comparative form indicates a higher property value for the subject or the object. That is that ‘big’ is covariant with size, as bigger things have a higher size value, and ‘small’ is contravariant with size.³ We also introduce a third concept of *absolute gradability*, which states that these objects are better described by these adjectives as they approach some ideal value. A common example of this is colours, where we may say that some object is redder than another if it is closer to some ideal value of red (e.g., RGB 0xff0000).

We have already mentioned absolute gradable adjectives above; it still appears here that absolute gradable adjectives are still gradable, which are (per definition) not. Absolute gradables are gradable and intersective!

While these concepts well handle the comparative usage of adjectives, the predicative and superlative usage of adjectives is complicated by three factors that we will outline below. We notice that gradable classes are not crisply defined as with the case of many intersective adjectives, that is that while we can clearly define all people in the world as ‘Belgian’ or ‘not Belgian’, by who holds a passport of Belgium, it is not easy split the world’s population into ‘tall’ and ‘not tall’ (This is known as *sorites* paradox (Bennett, 2006)). In fact, while it may be easy to say that someone with height 6’6” (198cm) is ‘tall’, it is not clear whether someone of height 6’ (182cm) is ‘tall’, also given that this height should be compared to average (different) height for a man, or a woman. As such, one frequently used way to deal with this class of vague adjectives (and nouns) is via fuzzy logic (Goguen, 1969; Zadeh, 1975; Zadeh, 1965; Dubois and Prade, 1988; Bennett, 2006). **John, you put here the following citations: Goguen 1969, Zadeh 1965, Zadeh 1975, Dubois 1988. Since they had not been added in the bib file, the compiler generated a “?”, I have added them. Check whether these are the papers that you were referring to. For your indicated reference ‘dubois1988’ i found DuboisPrade88en** Actually, these are from Francesca Secondly, we note that these class boundaries are non-monotonic, that is that with knowledge of more instances of the relative class we must revise our class boundaries. This is especially the case for the superlative as the discovery of a new tallest person in the world would remove the existing tallest person in the world from the class of tallest person in the world. This non-monotonicity also affects the class boundaries of the gradable class itself, for example in the 18th century the average height of a male was 5’5” (165cm).⁴ and as such a male of 6’ would have been considered clearly to be tall. As such, we can conclude that each instance added to our ontology must revise the class boundaries of a gradable class, hence leading to the fact that gradable adjectives are fundamentally non-monotonic. Finally, we must notice that gradability can only be understood relative to the class that we wish to grade, that is that while it is unclear as to whether 6’ is tall for a male, given the current average height of a female of about 5’4” (162cm) it is clear that 6’ is tall for a female.

²Note in many cases the property is quite abstract such as in ‘breakable’.

³The use of these terms is borrowed from type systems, and resembles the concept of ‘converse observables’ as mentioned by ((Bennett, 2006):42). As stated by the author, adjectives often come in pairs of polar opposites (e.g. *conv(tall) = short*, and both refer to the same observable (in this case *size*). Some observables analogously hold converse relationships with other observables (e.g. *conv(flexibility) = rigidity* or *conv(tallness) = shortness*).

⁴https://en.wikipedia.org/wiki/Human_height

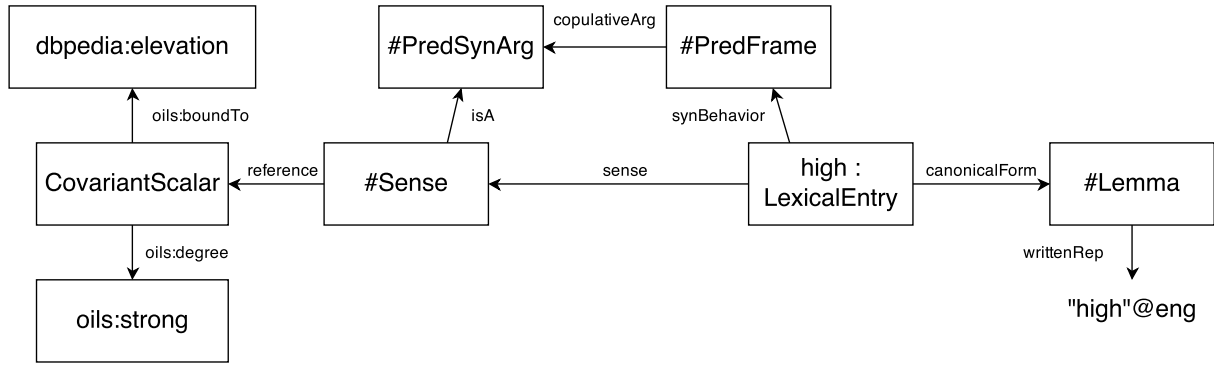


Figure 2: An example of the modelling of ‘high’ in *lemon*

So we can conclude that gradable adjectives are *fuzzy*, *non-monotonic* and *context-sensitive*, all of which are incompatible with the description logic used by OWL.

3.2.1 Pseudo-classes in lemonOILS

Currently there are only limited models for representing fuzzy logic in the context of the web (Zhao and Boley, 2008). In order to capture the properties of gradable adjectives, we introduce a new model, which we name *lemonOILS* (The *lemon* Ontology for the Interpretation of Lexical Semantics)⁵. This ontology introduces three new classes:

- *CovariantScalar*, indicating that the adjective is covariant with its bound property
- *ContravariantScalar*, indicating that the adjective is contravariant with its bound property
- *AbsoluteScalar*, indicating that the property represents similarity to an absolute value

In addition, the following properties are introduced to enable the description of gradable adjectives. Note, that all of these properties are typed as *annotation properties* in the OWL ontology, so that they do not interfere with the standard OWL reasoning.

- *boundTo* indicates the property that a scalar refers to (e.g., ‘size’ for ‘big’)
- *threshold* specifies a sensible minimal value for which the adjective can be said to hold
- *degree* is one of *weak*, *medium*, *strong* or *very strong*, corresponding to approximately 50%, 25%, 5% or 1% of all known individuals
- *comparator* indicates an object property that is equivalent to the comparison of the adjective (e.g., an object property *biggerThan* may be considered a comparator for the adjective class *big*)
- *measure* Indicates a unit that maybe used as a measure for this adjective, e.g., ‘John is 175 centimetres tall’, would use the ‘centimetre’ measurement.

Using such classes we can capture the semantics of gradable adjectives syntactically but not formally within an OWL model, as such we call such introduced classes *pseudo-classes*. An example of modelling an adjective such as ‘high’ is given below (depicted in Figure 2).

```
lexicon:high a lemon:LexicalEntry ;
  lemon:canonicalForm high:Lemma ;
  lemon:synBehavior high:PredFrame ;
  lemon:sense high:Sense .

high:Lemma lemon:writtenRep "high"@eng .

high:PredFrame lexinfo:copulativeArg high:PredArg .

high:Sense lemon:reference [
```

⁵<http://lemon-model.net/oils>

```

rdfs:subClassOf oils:CovariantScalar ;
oils:boundTo dbpedia:elevation ;
oils:degree oils:strong ] ;
lemon:isA high:PredArg .

```

As an example of a way in which it would be possible to interpret these annotations, we consider Markov Logic (Richardson and Domingos, 2006), which is an extension of first-order logic in which each clause is given a cost. The process of reasoning is thus transformed into an optimization problem of finding the extension, which minimizes the summed weight of all violated clauses. As such we can formulate a gradable adjective based on the number of known instances. For example, we can specify ‘big’ w.r.t. *size* for some class C as in 11.

$$\begin{aligned}
11. \quad & \forall x \in C, y \in C : size(x) > size(y) \rightarrow big_C(x) : \alpha \\
& \forall x \in C, y \in C : size(x) < size(y) \rightarrow \neg big_C(x) : \beta
\end{aligned}$$

In this way, the classification of an object into big or small can be defined as follows. For an individual $x \in C$ the property $big_C(x)$ holds if and only if:

$$|\{y \in C, size(y) > size(x)\}| \alpha < |\{y \in C, size(y) < size(x)\}| \beta$$

Where the values of α and β are related to the degree defined in the ontology.

We see that ‘big’ defined in this way has the three properties outlined above: it is non-monotonic (in that more individuals may change whether we consider an individual to be ‘big’ or not), it is fuzzy (given by the strength of the probability of the proposition $big_C(x)$), and it is context-sensitive (as whether an individual counts a big or not depends on the class C). Furthermore, our definition does not rely on defining big for a given class but instead is inferred from some known number of instances of this class. This eliminates the need to define a threshold for each individual class or even to define the predicate big_C on a per-class basis.

3.2.2 The supervaluation theory and SUMO

Another way to capture the meaning of these vague terms can be achieved by *supervaluation semantics*. Through supervaluation theory, the modeling or positioning of *sorites* vague concepts is grounded in a judgement or meaning that lies on arbitrary thresholds, but these thresholds are based on a number of *relevant objective measures* (Bennett, 2006).

A recent extension of the SUMO ontology (Niles and Pease, 2001, Suggested Upper Merged Ontology) has add default measurements for the Artifacts, Devices and Objects enlisted in the ontology (currently amounting to 300+). The compilation of defaultMeasurements in SUMO has been just conducted on observables. Given for instance an Artifact such as Book, the compilation of its default measurements would look like this:

```

;;Book
(defaultMinimumHeight Book (MeasureFn 10 Inch))
(defaultMaximumHeight Book (MeasureFn 11 Inch))
(defaultMinimumLength Book (MeasureFn 5.5 Inch))
(defaultMaximumLength Book (MeasureFn 7 Inch))
(defaultMinimumWidth Book (MeasureFn 1.2 Inch))
(defaultMaximumWidth Book (MeasureFn 5.5 Inch))

```

The example for Book shows that the default measurements for the observable reflect a *standard* kind of book, i.e., one of the most commonly known kinds of the same artifact. As for the case of Book, SUMO implies it to be a physical object with a certain length, height and width. A weakness here is that there is no systematic connection between the defaultMinimumHeight and Height, since properties like this have solely been defined in SUMO in terms of first-order logic, not in terms of default measurements. With *lemonOILS* we can add this information as follows:

```

sumo:Book oils:default [
  oils:defaultFor sumo:height ;
  oils:defaultMin "10in" ;
  oils:defaultMax "11in" ] .

```

Then if we understand a lexical entry ‘high’ as referring to a scalar covariant pseudo-class for `sumo:height`, it is possible to understand that a ‘high’ object exceeds the default minimum set established for the same and owns at the same time a value for ‘high’ which does not go beyond the established default maximum. A further weakness of this approach is captured by the following example:

12. Avery Johnson is a short basketball player.

Here, we see the difficulty in interpreting the sentence as Avery Johnson is in fact of average height (5’10”) but for the class of basketball players he is unusually short. While SUMO has some very specific listings of subsets for the same `Artifact`⁶ and SUMO does not provide a well-structured subset net for `Person` as such we must introduce default values for every possible subclass of person. A way to solve this bottleneck could be to introduce, for every possible subclass of person, default values, as well as to introduce default values for the same `Artifact` in conjunction with a predicate or adjective. Nevertheless, this can potentially lead to an explosion in the amount of information we need to encode in the ontology.

3.3 Operator adjectives

To clarify why we introduce operator and relational adjectives at this point, we might want to tell that these kinds of adjectives are not gradable, but they are usually related to their ‘classificatory’ or ‘taxonomic’ meaning (McNally and Boleda, 2004), which should make them particularly suitable for being modeled in ontologies (although we present here some bottlenecks as well). Both operator and relational adjectives are characterized by a not-postnominal position (this is why I don’t think ‘known’ below is an ok example for relational adjective [we can in fact say ‘he is known’ out of ‘he is a known actor’]).

13. Clinton is a male tall former president
14. \models Clinton is president.
15. \models Clinton is male (male \Rightarrow intersective gradable adjective)
16. \models Clinton is tall (tall \Rightarrow subsective gradable adjective)
17. \emptyset Clinton is former. (Since this example was the same as (2a,b) above, I’ve changed 2(a)(b) into ‘Mary is a technical engineer \Rightarrow *Mary is technical’)

Relational adjectives and operator adjectives are not exclusively attributive. But, as they require some secondary argument (class or individual resp.) they can only be understood with this second argument, although that argument may be omitted, e.g., he is a ‘known person’ means he is known to somebody (e.g., the police)

Operator adjectives are those that combine with a noun to modify the meaning of the noun itself. There are two primary issues with the understanding of the adjective in this manner. Firstly, the reference of the lexical item does not generally refer to an existing item in the ontology, but rather is novel and productive. Secondly, the compositional nature of adjective-noun compounds is no longer simple as in the cases of intersective and gradable adjectives. This means that in order to understand a concept such as a ‘fake gun’ we must first derive a class of `FakeGuns` from the class of `Guns`. Thus the modified noun phrase must be an argument of the operator adjective. To this extent we claim that it is not generally possible to represent the meaning of an operator adjective within the context of an OWL ontology. Instead, following Bankston (Bankston, 2003), we claim that the reference of an operator adjective must be a higher order predicate. If we assume that there are operators of the form of a function, then the argument of an operator is the attributed noun phrase. As such we introduce a frame *operator attributive*, that has one argument which is the noun. The interpretation of a phrase like 18a is like 18b.

18. (a) Obama is the president. Clinton is a former president.
 (b) $President(Obama) \sqcap [former(President)](Clinton)$

⁶For example, the subsets for `Car` are `CrewDormCar`, `GalleryCar` and more.

If we understand *President* as *Artifact* or *Human* that has the role of leading a country (and therefore being true for *Obama* and *Clinton*), then *former* creates a new class in the ontology (i.e. *FormerPresident*) which selects people who were presidents. An alternative interpretation would be that ‘former’ modifies the role of *Clinton*, however this raises the question of why the role property of presidents would be selected, and as such is also unsatisfactory. As there is at the moment no agreed representation for such an operator in an ontology, it underlines the fact that more sophisticated ontology representations than first-order logic are required to understand natural language in general.

3.4 Relational adjectives

Relational adjectives⁷ are those that require a second argument, such as ‘known’, which can only be understood as being ‘known’ to some person, in comparison to ‘famous’. As such the modelling of relational adjectives is similar to we can model ‘known’ with the frame ‘*X* is known to *Y*’ and reference foaf:knows as: **Another example, apart from ‘known’, (that you can use in postnominal position, and therefore does not really represent relational adjective) could be: ‘Jemma is a single parent’ ⇒ ‘Jemma is a parent’, but ∅ ‘Jemma is single’ ‘single’ is an operator here, it adds the restriction (‘there can be only one’) to the class of Parent, ‘Jemma is single’ is another sense of ‘single’**

```
lexicon:known a lemon:LexicalEntry ;
  lemon:canonicalForm known:Lemma ;
  lemon:sense known:Sense ;
  lemon:synBehavior known:Frame .

known:Lemma lemon:writtenRep "known"@eng .

known:Frame lexinfo:attributeArg known:Subject ;
  lexinfo:prepositionalObject known:Object .

known:Sense lemon:reference foaf:knows ;
  lemon:subjOfProp known:Subject ;
  lemon:objOfProp known:Object .

known:Object lemon:marker lexicon:to .
```

Yeah, perhaps we should find another name? Any suggestions?

4 Adjectives in question answering

Of the 150 training and test questions of the QALD-4 benchmark⁸ for question answering over linked data, 67 contain adjectives.

Some of the occurring adjectives do not have a semantic contribution, or at least none that is separable from the noun, as exemplified in the noun phrases in 19 and 20⁹.

19. (a) $\llbracket \text{official website} \rrbracket = \text{dbo:website}$
 (b) $\llbracket \text{national anthem} \rrbracket = \text{dbo:anthem}$
20. (a) $\llbracket \text{official languages} \rrbracket = \text{dbo:officialLanguages}$
 (b) $\llbracket \text{military conflicts} \rrbracket = \text{dbo:battle}$

Otherwise, the most common kinds of adjectives among them are intersective and gradable adjectives.

All intersective adjectives denote restriction classes that are not explicitly named in DBpedia, in correspondence with the modelling proposed in Section 3.1 above, for example:

21. (a) $\llbracket \text{Danish} \rrbracket = \exists \text{dbo:country.res:Denmark}$
 (b) $\llbracket \text{female} \rrbracket = \exists \text{dbo:gender.res:Female}$
 (c) $\llbracket \text{Methodist} \rrbracket = \exists \text{dbo:religion.res:Methodism}$

⁷Our definition of relational here is based on the idea of relational nouns (De Bruin and Scha, 1988) as a word that requires an argument. It is important to specify that relational adjectives can also have another definition in literature (Morzycki, 2013a).

⁸<http://www.sc.cit-ec.uni-bielefeld.de/qald/>

⁹ $\llbracket \cdot \rrbracket$ stands for ‘denotes’ and the prefixes *dbo* and *res* abbreviate the DBpedia namespaces <http://dbpedia.org/ontology/> and <http://dbpedia.org/resource/>, respectively

In some cases these intersectives have a context-dependent and highly ontology-specific meaning, often tightly interwoven with the meaning of the noun, as in the following examples:

22. (a) $\llbracket \text{first president of the United States} \rrbracket = \exists \text{dbo:office} . \text{'1st President of the United States'}$
 (b) $\llbracket \text{first season} \rrbracket = \exists \text{dbo:seasonNumber} . 1$

All gradable adjectives that occur in the QALD-4 question set can be captured in terms of *lemonOILS* as *CovariantScalar* (e.g. ‘high’) or *ContravariantScalar* (e.g. ‘young’), cf. Section 3.2 above, bound to a DBpedia datatype property (e.g. *elevation* or *birthDate*). The positive form of those adjectives only occurs in ‘how’ questions, denoting the property they are bound to, for example:

23. (a) $\llbracket \text{deep} \rrbracket = \text{dbo:depth}$ in ‘How deep is Lake Placid?’
 (b) $\llbracket \text{tall} \rrbracket = \text{dbo:height}$ in ‘How tall is Michael Jordan?’

The comparative form denotes the property they are bound to together with an aggregation operation, usually a filter invoking a term of comparison that depends on whether the adjective is covariant or contravariant.

24. (a) $\llbracket \text{Which mountains are higher than the Nanga Parbat?} \rrbracket =$

```
SELECT DISTINCT ?uri WHERE {
  res:Nanga_Parbat dbo:elevation ?x .
  ?uri rdf:type dbo:Mountain .
  ?uri dbo:elevation ?y .
  FILTER (?y > ?x)
}
```

Finally, the superlative form denotes the property they are bound to together with an aggregation operation, usually an ordering with a cut-off of all results except the first one, as exemplified in 25. In some cases the superlative property is already encoded in the ontology, e.g., *dbo:highestPlace*.¹⁰

25. $\llbracket \text{What is the longest river?} \rrbracket =$

```
SELECT DISTINCT ?uri WHERE {
  ?uri rdf:type dbo:River .
  ?uri dbo:length ?l .
} ORDER BY DESC(?l) OFFSET 0 LIMIT 1
```

There are some instances of operator adjectives, e.g. ‘former’, which is used in a context where it does not have a specific semantic contribution:

26. $\llbracket \text{the former Dutch queen Juliana} \rrbracket = \text{res:Juliana}$

Finally, there is only one instance of an operator adjective:

27. $\llbracket \text{professional surfer} \rrbracket = \exists \text{dbo:occupation} . \text{res:Surfing}$

5 Related work

The categorization of adjectives in terms of formal semantics goes back to Montague (1970) and Vendler (1968), however one of the most significant attempts to assign a formal meaning was carried out in the Mikrokosmos project (Raskin and Nirenburg, 1995). This was one of the first works to treat the case of a micro-theory of adjectives, in which the results were “machine-tractable”, in that they could be formally defined by a computer. The applications of this were limited however and no formal logic was attached to the semantic representations. Nevertheless much of the modeling resembles ours. In particular, scalar adjectives in the Mikrokosmos project are modeled by association with an attribute and a range, e.g., ‘big’ is described as being > 0.75 (i.e., 75% of all known instances) on the

¹⁰The prefix *dbp* abbreviates the DBpedia namespace <http://dbpedia.org/property/>

size-attribute. These classifications do not however clearly separate meaning and syntax and also require a separate modeling of comparatives and class-specific meanings for many adjectives.

Amoia and Garden (2006) handled the problem of adjectives in the context of textual entailment. They analyzed 15 classes that show the subtle interaction between the semantic class (e.g., ‘privative’) and the issues of attributive/predicative use and gradability. Abdullah and Frost (2005) tackles the privative nature of adjectives by arguing that the adjectives modify the set themselves, in a manner that is naturally second-order. Similarly, Partee (2003) proposed a limited second-order model by means of an ‘head primary principle’ requiring that adjectives are interpreted within their context. Bankston’s analysis (2003) however shows that the fundamental nature of many adjectives is higher-order, and provides a very sophisticated formal representation framework for this syntactic class. A more thorough discussion of non-gradable, non-intersective adjectives is given by Morzycki (2013a).

The Generative Lexicon (Pustejovsky, 1991) provides another approach to the representation of semantics, and the case of adjectives has also been considered in this context Bouillion (1999) consider the case of the French adjective ‘vieux’ (‘old’), which he interprets as selecting two different elements in the event structure of an attributed noun, that is whether the state, e.g., being a ‘mayor’ for ‘mayor’, is considered old or the individual itself. In this way, the introduction of two senses for ‘vieux’ is avoided, however it remains unclear if such reasoning introduces more complexity than the extra senses. Furthermore, In his analysis of adjectives, Larson (1998) suggests that many adjectives denote properties of *events*, rather than of simple heads or nouns (which does not fall very far from the statement, made above, that relational adjectives denote properties of kinds). Pustejovsky (1992; 1991) and (Lenci and others, 2000) state that lexical and semantic decomposition can be achieved *generatively*, assigning to each lexical item a specific *qualia structure*. For instance, in an expression like:

28. The round, heavy, wooden, inlaid magnifying glass

- ‘round’ represents the *Formal* role (giving indications of shape and dimensionality)
- ‘heavy’ and ‘wooden’ is the *Constitutive* role and indicates the relation between the object and its parts (e. g. by specifying weight, material, parts and components)
- ‘inlaid’ is the *Agentive* role of the lexical item, denoting the factors that have been involved in the generation of the objects, such as creator, artifact, natural kind, and causal chain
- ‘magnifying’ describes the *Telic* role of ‘glass’, since it shows its purpose and function

Finally, Peters and Peters (2000) provide one of the few other practical reports on modelling adjectives with ontologies, in the context of the SIMPLE lexica. This work is primarily focussed on the categorization of by means of intensional and extensional properties, rather than due to their logical modelling.

6 Conclusion

In this paper we have presented a method for modelling adjectives with the ontology-lexicon model, *lemon*. In particular, we found that adjectives frequently go beyond the first-order logic model used by OWL, but instead require models that are non-monotonic, fuzzy and second-order. As such, we conclude that more sophisticated semantic models are required to represent the semantics of such words. However the separation of syntax and semantics remains a robust model, which can easily be adapted to the task of representing adjectives. As a final note, we consider the claims that not all languages have adjectives¹¹ however we contend that the underlying semantics of the words we discuss here is representable in all nearly languages. Based on our analysis of realistic questions as applied in QALD, we furthermore believe that this model should be applicable to a range of domains and languages with little issue, even if further validation is naturally necessary.

References

Nabil Abdullah and Richard A Frost. 2005. Adjectives: A uniform semantic approach. In *Advances in Artificial Intelligence*, pages 330–341. Springer.

¹¹<http://linguistlist.org/issues/4/4-442.html>

- Marilisa Amoia and Claire Gardent. 2006. Adjective based inference. In *Proceedings of the Workshop KRAQ'06 on Knowledge and Reasoning for Language Processing*, pages 20–27. Association for Computational Linguistics.
- Paul Bankston. 2003. Modeling nonintersective adjectives using operator logics. *The Review of Modern Logic*, 9(1-2):9–28.
- Brandon Bennett. 2006. A theory of vague adjectives grounded in relevant observables. In John Mylopoulos Patrick Doherty and Christopher A. Welty, editors, *Proceedings of the Tenth International Conference on Principles of Knowledge Representation and Reasoning*, pages 36–45. AAAI Press.
- Pierrette Bouillon and Evelyne Viegas. 1999. The description of adjectives for natural language processing: Theoretical and applied perspectives. In *Proceedings of Description des Adjectifs pour les Traitements Informatiques. Traitement Automatique des Langues Naturelles*. Citeseer.
- Pierrette Bouillon. 1999. The adjective “vieux”: The point of view of “generative lexicon”. In *Breadth and depth of semantic lexicons*, pages 147–166. Springer.
- Paul Buitelaar, 2010. *Ontology-based Semantic Lexicons: Mapping between Terms and Object Descriptions*, pages 212–223. Cambridge University Press.
- Philipp Cimiano, Paul Buitelaar, John M^cCrae, and Michael Sintek. 2011. Lexinfo: A declarative model for the lexicon-ontology interface. *Web Semantics: Science, Services and Agents on the World Wide Web*, 9(1):29–51.
- Philipp Cimiano, Janna Lüker, David Nagel, and Christina Unger. 2013. Exploiting ontology lexica for generating natural language texts from rdf data. In *Proceedings of the 14th European Workshop on Natural Language Generation*, pages 10–19.
- Jos De Bruin and Remko Scha. 1988. The interpretation of relational nouns. In *Proceedings of the 26th annual meeting on Association for Computational Linguistics*, pages 25–32. Association for Computational Linguistics.
- Frank Van Harmelen Deborah L. M^cGuinness et al. 2004. Owl web ontology language overview. *W3C recommendation*, 10(2004-03):10.
- Didier Dubois and Henri Prade. 1988. *Possibility theory*. Plenum Press, New York.
- Joseph H. Goguen. 1969. The logic of inexact concepts. *Synthese*, 19:325–373.
- Christopher Kennedy and Louise McNally. 1999. Deriving the scalar structure of deverbal adjectives. *Catalan Working Papers in Linguistics*, 7:125–139.
- Christopher Kennedy. 2007. Vagueness and grammar: the semantics of relative and absolute gradable adjectives. *Linguistics and philosophy*, 30:1–45.
- Richard K. Larson. 1998. Events and modification in nominals. In Devon Strolovitch and Aaron Lawson, editors, *Proceedings from Semantics and Linguistic Theory (SALT) VIII*, pages 145–168. CLC Publications, Itaca, New York.
- Alessandro Lenci et al. 2000. Simple work package 2, linguistic specifications, deliverable d2.1.
- Vanessa Lopez, Christina Unger, Philipp Cimiano, and Enrico Motta. 2013. Evaluating question answering over linked data. *Web Semantics: Science, Services and Agents on the World Wide Web*, 21:3–13.
- Louise McNally and Gemma Boleda. 2004. Relational adjectives as properties of kinds. *Empirical issues in formal syntax and semantics*, 5:179–196.
- Richard Montague. 1970. English as a formal language. In Bruno Visentini et al, editor, *Linguaggi nella società e nella tecnica*, pages 189–224. Milan: Edizioni di Comunità.
- Marcin Morzycki. 2013a. The lexical semantics of adjectives: More than just scales. Ms., Michigan State University. Draft of a chapter in *Modification*, a book in preparation for the Cambridge University Press series *Key Topics in Semantics and Pragmatics*.
- Marcin Morzycki. 2013b. *Modification*. Cambridge University Press.
- John P. M^cCrae and Christina Unger. 2014. Design patterns for the ontology-lexicon interface. In Paul Buitelaar and Philipp Cimiano, editors, *Towards the Multilingual Semantic Web: Principles, Methods and Applications*. Springer.
- John M^cCrae, Guadalupe Aguado-de Cea, Paul Buitelaar, Philipp Cimiano, Thierry Declerck, Asunción Gómez-Pérez, Jorge Gracia, Laura Hollink, Elena Montiel-Ponsoda, Dennis Spohr, et al. 2012. Interchanging lexical resources on the semantic web. *Language Resources and Evaluation*, 46(4):701–719.
- Ian Niles and Adam Pease. 2001. Towards a standard upper ontology.
- Barbara H Partee. 2003. Are there privative adjectives. In *Conference on the Philosophy of Terry Parsons, University of Massachusetts, Amherst*.
- Ivonne Peters and Wim Peters. 2000. The treatment of adjectives in simple: Theoretical observations. In *LREC*.

- James Pustejovsky. 1991. The generative lexicon. *Computational linguistics*, 17(4):409–441.
- James Pustejovsky. 1992. The syntax of event structure. In Bett Levin and Steven Pinker, editors, *Lexical & Conceptual Semantics*, pages 47–83. Oxford: Blackwell.
- Victor Raskin and Sergei Nirenburg. 1995. Lexical semantics of adjectives. *New Mexico State University, Computing Research Laboratory Technical Report, MCCS-95-288*.
- Matthew Richardson and Pedro Domingos. 2006. Markov logic networks. *Machine learning*, 62(1-2):107–136.
- Alexandra Teodorescu. 2006. Adjective ordering restrictions revisited. In *Proceedings of the 25th West Coast Conference on Formal Linguistics*, pages 399–407. Citeseer.
- Christina Unger and Philipp Cimiano. 2011. Pythia: Compositional meaning construction for ontology-based question answering on the semantic web. In Rafael Munoz, editor, *Natural Language Processing and Information Systems: 16th International Conference on Applications of Natural Language to Information Systems, NLDB 2011, Alicante, Spain, June 28-30, 2011. Proceedings*, volume 6716, pages 153–160. Springer.
- Zeno Vendler. 1968. *Adjectives and nominalizations*. Number 5 in Papers on formal linguistics. Mouton.
- Lofti A. Zadeh. 1965. Fuzzy sets. *Information and Control*, 8:338–353.
- Lofti A. Zadeh. 1975. The concept of linguistic variable and its application to approximate reasoning. *Information Sciences*, 8:199–249.
- Jidi Zhao and Harold Boley. 2008. Uncertainty treatment in the rule interchange format: From encoding to extension. In *URSW*.